

```
In [1]: import math
import numpy as np
import pandas as pd
from matplotlib import pyplot as plt
from mpl_toolkits.mplot3d import Axes3D
from sklearn.model_selection import train_test_split
from sklearn.model_selection import StratifiedKFold
from sklearn.svm import SVC
from sklearn.metrics import average_precision_score
```

Reading in the data

```
In [2]: sales_data = pd.read_csv("wholesale-customers.csv")
```

```
In [3]: sales_data.shape
```

```
Out[3]: (440, 8)
```

```
In [4]: sales_data.head(5)
```

```
Out[4]:
```

	Channel	Region	Fresh	Milk	Grocery	Frozen	Detergents_Paper	Delicassen
0	2	3	12669	9656	7561	214	2674	1338
1	2	3	7057	9810	9568	1762	3293	1776
2	2	3	6353	8808	7684	2405	3516	7844
3	1	3	13265	1196	4221	6404	507	1788
4	2	3	22615	5410	7198	3915	1777	5185

```
In [5]: sales_data.tail(5)
```

```
Out[5]:
```

	Channel	Region	Fresh	Milk	Grocery	Frozen	Detergents_Paper	Delicassen
435	1	3	29703	12051	16027	13135	182	2204
436	1	3	39228	1431	764	4510	93	2346
437	2	3	14531	15488	30243	437	14841	1867
438	1	3	10290	1981	2232	1038	168	2125
439	1	3	2787	1698	2510	65	477	52

Train/Test Split

```
In [6]: X = sales_data.iloc[:, 2:]
```

```
In [7]: x.head(5)
```

```
Out[7]:
```

	Fresh	Milk	Grocery	Frozen	Detergents_Paper	Delicassen
0	12669	9656	7561	214	2674	1338
1	7057	9810	9568	1762	3293	1776
2	6353	8808	7684	2405	3516	7844
3	13265	1196	4221	6404	507	1788
4	22615	5410	7198	3915	1777	5185

```
In [8]: x.tail(5)
```

```
Out[8]:
```

	Fresh	Milk	Grocery	Frozen	Detergents_Paper	Delicassen
435	29703	12051	16027	13135	182	2204
436	39228	1431	764	4510	93	2346
437	14531	15488	30243	437	14841	1867
438	10290	1981	2232	1038	168	2125
439	2787	1698	2510	65	477	52

```
In [9]: y = sales_data.iloc[:, [0]]
```

```
In [10]: y.head(5)
```

```
Out[10]:
```

	Channel
0	2
1	2
2	2
3	1
4	2

```
In [11]: y.tail(5)
```

```
Out[11]:
```

	Channel
435	1
436	1
437	2
438	1
439	1

```
In [12]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, ra
```

Model Development

Linear Support Vector Machine

Perform stratified k-fold cross validation

```
In [13]: avg_precision_recall_scores = [0, 0, 0, 0, 0]  
c_values = [0.0001, 0.001, 0.01, 0.1, 1.0]  
num_iterations = 0
```

```
In [19]: kf = StratifiedKFold(n_splits = 2)

for train_index, test_index in kf.split(X,y):
    X_train_indexed, X_test_indexed = X.iloc[train_index, :], X.iloc[test_index, :]
    y_train_indexed, y_test_indexed = y.iloc[train_index, :], y.iloc[test_index, :]

    for i in range(5):
        c_value = c_values[i]
        linear_svm_model = SVC(C = c_value, kernel = "linear")
        linear_svm_model.fit(X_train_indexed, y_train_indexed)
        y_score = linear_svm_model.decision_function(X_test_indexed)
        avg_precision_recall_score = average_precision_score(y_test_indexed, y_score)
        avg_precision_recall_scores[i] += avg_precision_recall_score
        num_iterations += 1
```

/Users/ryanlim/anaconda3/lib/python3.7/site-packages/sklearn/utils/validation.py:761: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

y = column_or_1d(y, warn=True)

/Users/ryanlim/anaconda3/lib/python3.7/site-packages/sklearn/utils/validation.py:761: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

y = column_or_1d(y, warn=True)

/Users/ryanlim/anaconda3/lib/python3.7/site-packages/sklearn/utils/validation.py:761: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

y = column_or_1d(y, warn=True)

/Users/ryanlim/anaconda3/lib/python3.7/site-packages/sklearn/utils/validation.py:761: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

y = column_or_1d(y, warn=True)

/Users/ryanlim/anaconda3/lib/python3.7/site-packages/sklearn/utils/validation.py:761: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

y = column_or_1d(y, warn=True)

/Users/ryanlim/anaconda3/lib/python3.7/site-packages/sklearn/utils/validation.py:761: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

y = column_or_1d(y, warn=True)

/Users/ryanlim/anaconda3/lib/python3.7/site-packages/sklearn/utils/validation.py:761: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

y = column_or_1d(y, warn=True)

/Users/ryanlim/anaconda3/lib/python3.7/site-packages/sklearn/utils/validation.py:761: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

y = column_or_1d(y, warn=True)

/Users/ryanlim/anaconda3/lib/python3.7/site-packages/sklearn/utils/validation.py:761: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().


d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

```
y = column_or_1d(y, warn=True)
/Users/ryanlim/anaconda3/lib/python3.7/site-packages/sklearn/utils/validation.py:761: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples, ), for example using ravel().
y = column_or_1d(y, warn=True)
```

```
In [20]: averaged_missed_precision_recall_scores = [(1 - avg_precision_recall_score)
```

Create a plot of the c (tuning parameter) and average missed precision-recall values

```
In [21]: plt.plot(c_values, averaged_missed_precision_recall_scores )
plt.xlabel("C (Tuning Parameter) Value")
plt.ylabel("Avg Missed Precision-Recall Value")
plt.show()
```



Find the optimal c (tuning parameter) value and use it to build a linear support vector machine model.

```
In [22]: min_avg_missed_precision_recall_score = min(averaged_missed_precision_recall_scores)
min_error_index = averaged_missed_precision_recall_scores.index(min_avg_missed_precision_recall_score)
c_value = c_values[min_error_index]
```

```
In [23]: linear_svm_model = SVC(C = c_value, kernel = "linear")
linear_svm_model.fit(X_train, y_train)
y_score = linear_svm_model.decision_function(X_test)
avg_precision_recall_score = average_precision_score(y_test, y_score)
avg_precision_recall_scores[i] += avg_precision_recall_score
```

/Users/ryanlim/anaconda3/lib/python3.7/site-packages/sklearn/utils/validation.py:761: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

```
y = column_or_1d(y, warn=True)
```

Specify the linear support vector machine model coefficients and intercept

```
In [24]: coefficients = pd.DataFrame({"Features":X.columns, "Coefficients":np.transpose(coefficients.head(8))
```

Out[24]:

	Features	Coefficients
0	Fresh	-0.001861
1	Milk	0.002976
2	Grocery	0.001329
3	Frozen	-0.009764
4	Detergents_Paper	0.024569
5	Delicassen	0.004162

```
In [25]: intercept = pd.DataFrame({"Intercept": linear_svm_model.intercept_})
intercept.head()
```

Out[25]:

	Intercept
0	-80.197733

Specify the minimum cross validation error

```
In [26]: print("The minimum average cross validation missed precision-recall score is: ")
```

```
The minimum average cross validation missed precision-recall score is: 0.004947229711484813
```

Support Vector Machine With Polynomial Kernel

Perform stratified k-fold cross validation

```
In [27]: avg_precision_recall_scores = [0, 0, 0, 0, 0]
c_values = [0.0001, 0.001, 0.01, 0.1, 1.0]
num_iterations = 0
```

```
In [29]: kf = StratifiedKFold(n_splits = 2)

for train_index, test_index in kf.split(X,y):
    X_train_indexed, X_test_indexed = X.iloc[train_index, :], X.iloc[test_index, :]
    y_train_indexed, y_test_indexed = y.iloc[train_index, :], y.iloc[test_index, :]

    for i in range(5):
        c_value = c_values[i]
        polynomial_kernel_svm_model = SVC(C = c_value, kernel = "poly")
        polynomial_kernel_svm_model.fit(X_train_indexed, y_train_indexed)
        y_score = polynomial_kernel_svm_model.decision_function(X_test_indexed)
        avg_precision_recall_score = average_precision_score(y_test_indexed, y_score)
        avg_precision_recall_scores[i] += avg_precision_recall_score
    num_iterations += 1
```

/Users/ryanlim/anaconda3/lib/python3.7/site-packages/sklearn/utils/validation.py:761: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

y = column_or_1d(y, warn=True)

/Users/ryanlim/anaconda3/lib/python3.7/site-packages/sklearn/svm/base.py:196: FutureWarning: The default value of gamma will change from 'auto' to 'scale' in version 0.22 to account better for unscaled features. Set gamma explicitly to 'auto' or 'scale' to avoid this warning.

"avoid this warning.", FutureWarning)

/Users/ryanlim/anaconda3/lib/python3.7/site-packages/sklearn/utils/validation.py:761: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

y = column_or_1d(y, warn=True)

/Users/ryanlim/anaconda3/lib/python3.7/site-packages/sklearn/svm/base.py:196: FutureWarning: The default value of gamma will change from 'auto' to 'scale' in version 0.22 to account better for unscaled features. Set gamma explicitly to 'auto' or 'scale' to avoid this warning.

"avoid this warning.", FutureWarning)

/Users/ryanlim/anaconda3/lib/python3.7/site-packages/sklearn/utils/validation.py:761: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

y = column_or_1d(y, warn=True)

/Users/ryanlim/anaconda3/lib/python3.7/site-packages/sklearn/svm/base.py:196: FutureWarning: The default value of gamma will change from 'auto' to 'scale' in version 0.22 to account better for unscaled features. Set gamma explicitly to 'auto' or 'scale' to avoid this warning.

"avoid this warning.", FutureWarning)

/Users/ryanlim/anaconda3/lib/python3.7/site-packages/sklearn/utils/validation.py:761: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

y = column_or_1d(y, warn=True)

/Users/ryanlim/anaconda3/lib/python3.7/site-packages/sklearn/svm/base.py:196: FutureWarning: The default value of gamma will change from 'auto' to 'scale' in version 0.22 to account better for unscaled features. Set gamma explicitly to 'auto' or 'scale' to avoid this warning.

"avoid this warning.", FutureWarning)

/Users/ryanlim/anaconda3/lib/python3.7/site-packages/sklearn/utils/validation.py:761: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

```

d array was expected. Please change the shape of y to (n_samples, ), for
example using ravel().
y = column_or_1d(y, warn=True)
/Users/ryanlim/anaconda3/lib/python3.7/site-packages/sklearn/svm/base.py:
196: FutureWarning: The default value of gamma will change from 'auto' to
'scale' in version 0.22 to account better for unscaled features. Set gamm
a explicitly to 'auto' or 'scale' to avoid this warning.
"avoid this warning.", FutureWarning)
/Users/ryanlim/anaconda3/lib/python3.7/site-packages/sklearn/utils/valida
tion.py:761: DataConversionWarning: A column-vector y was passed when a 1
d array was expected. Please change the shape of y to (n_samples, ), for
example using ravel().
y = column_or_1d(y, warn=True)
/Users/ryanlim/anaconda3/lib/python3.7/site-packages/sklearn/svm/base.py:
196: FutureWarning: The default value of gamma will change from 'auto' to
'scale' in version 0.22 to account better for unscaled features. Set gamm
a explicitly to 'auto' or 'scale' to avoid this warning.
"avoid this warning.", FutureWarning)
/Users/ryanlim/anaconda3/lib/python3.7/site-packages/sklearn/utils/valida
tion.py:761: DataConversionWarning: A column-vector y was passed when a 1
d array was expected. Please change the shape of y to (n_samples, ), for
example using ravel().
y = column_or_1d(y, warn=True)
/Users/ryanlim/anaconda3/lib/python3.7/site-packages/sklearn/svm/base.py:
196: FutureWarning: The default value of gamma will change from 'auto' to
'scale' in version 0.22 to account better for unscaled features. Set gamm
a explicitly to 'auto' or 'scale' to avoid this warning.
"avoid this warning.", FutureWarning)
/Users/ryanlim/anaconda3/lib/python3.7/site-packages/sklearn/utils/valida
tion.py:761: DataConversionWarning: A column-vector y was passed when a 1
d array was expected. Please change the shape of y to (n_samples, ), for
example using ravel().
y = column_or_1d(y, warn=True)
/Users/ryanlim/anaconda3/lib/python3.7/site-packages/sklearn/svm/base.py:
196: FutureWarning: The default value of gamma will change from 'auto' to
'scale' in version 0.22 to account better for unscaled features. Set gamm
a explicitly to 'auto' or 'scale' to avoid this warning.
"avoid this warning.", FutureWarning)
/Users/ryanlim/anaconda3/lib/python3.7/site-packages/sklearn/utils/valida
tion.py:761: DataConversionWarning: A column-vector y was passed when a 1
d array was expected. Please change the shape of y to (n_samples, ), for
example using ravel().
y = column_or_1d(y, warn=True)
/Users/ryanlim/anaconda3/lib/python3.7/site-packages/sklearn/svm/base.py:
196: FutureWarning: The default value of gamma will change from 'auto' to
'scale' in version 0.22 to account better for unscaled features. Set gamm

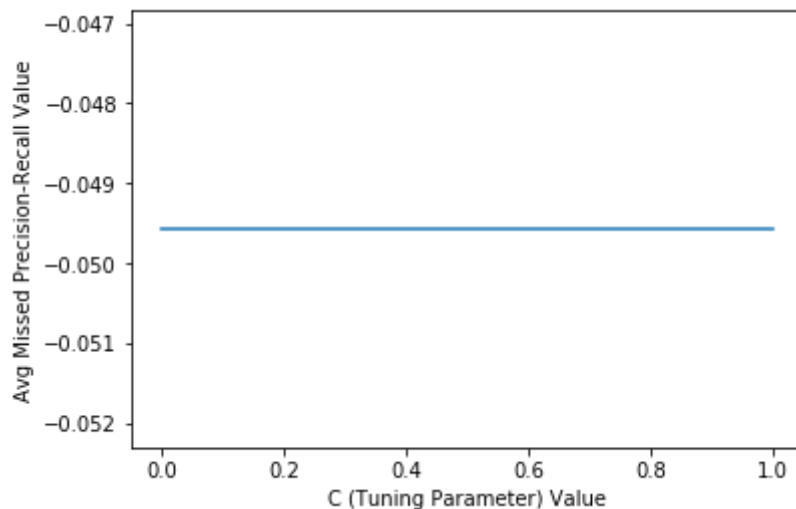
```


a explicitly to 'auto' or 'scale' to avoid this warning.
 "avoid this warning.", FutureWarning)

```
In [31]: averaged_missed_precision_recall_scores = [(1 - avg_precision_recall_score)
```

Create a plot of the c (tuning parameter) and average missed precision-recall values

```
In [32]: plt.plot(c_values, averaged_missed_precision_recall_scores )
plt.xlabel("C (Tuning Parameter) Value")
plt.ylabel("Avg Missed Precision-Recall Value")
plt.show()
```



Specify the minimum cross_validation error

```
In [33]: min_avg_missed_precision_recall_score = min(averaged_missed_precision_recall_scores)
print(min_avg_missed_precision_recall_score)
```

```
-0.049571645366501674
```

Support Vector Machine With Gaussian Kernel

Perform stratified k-fold cross validation

```
In [34]: avg_precision_recall_scores = [0, 0, 0, 0, 0]
c_values = [0.0001, 0.001, 0.01, 0.1, 1.0]
num_iterations = 0
```

```
In [35]: kf = StratifiedKFold(n_splits = 2)

for train_index, test_index in kf.split(X,y):
    X_train_indexed, X_test_indexed = X.iloc[train_index, :], X.iloc[test_index, :]
    y_train_indexed, y_test_indexed = y.iloc[train_index, :], y.iloc[test_index, :]

    for i in range(5):
        c_value = c_values[i]
        gaussian_kernel_svm_model = SVC(C = c_value, kernel = "rbf", gamma=
        gaussian_kernel_svm_model.fit(X_train_indexed, y_train_indexed)
        y_score = gaussian_kernel_svm_model.decision_function(X_test_indexed)
        avg_precision_recall_score = average_precision_score(y_test_indexed, y_score)
        avg_precision_recall_scores[i] += avg_precision_recall_score
    num_iterations += 1
```

/Users/ryanlim/anaconda3/lib/python3.7/site-packages/sklearn/utils/validation.py:761: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

y = column_or_1d(y, warn=True)

/Users/ryanlim/anaconda3/lib/python3.7/site-packages/sklearn/utils/validation.py:761: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

y = column_or_1d(y, warn=True)

/Users/ryanlim/anaconda3/lib/python3.7/site-packages/sklearn/utils/validation.py:761: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

y = column_or_1d(y, warn=True)

/Users/ryanlim/anaconda3/lib/python3.7/site-packages/sklearn/utils/validation.py:761: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

y = column_or_1d(y, warn=True)

/Users/ryanlim/anaconda3/lib/python3.7/site-packages/sklearn/utils/validation.py:761: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

y = column_or_1d(y, warn=True)

/Users/ryanlim/anaconda3/lib/python3.7/site-packages/sklearn/utils/validation.py:761: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

y = column_or_1d(y, warn=True)

/Users/ryanlim/anaconda3/lib/python3.7/site-packages/sklearn/utils/validation.py:761: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

y = column_or_1d(y, warn=True)

/Users/ryanlim/anaconda3/lib/python3.7/site-packages/sklearn/utils/validation.py:761: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

y = column_or_1d(y, warn=True)

/Users/ryanlim/anaconda3/lib/python3.7/site-packages/sklearn/utils/validation.py:761: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

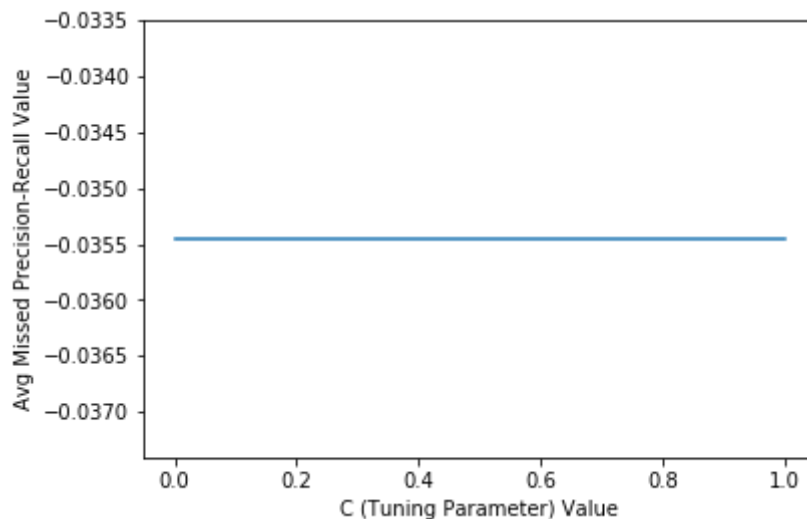
d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

```
y = column_or_1d(y, warn=True)
/Users/ryanlim/anaconda3/lib/python3.7/site-packages/sklearn/utils/validation.py:761: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples, ), for example using ravel().
y = column_or_1d(y, warn=True)
```

```
In [36]: averaged_missed_precision_recall_scores = [(1 - avg_precision_recall_score)
```

Create a plot of the c (tuning parameter) and average missed precision-recall values

```
In [37]: plt.plot(c_values, averaged_missed_precision_recall_scores)
plt.xlabel("C (Tuning Parameter) Value")
plt.ylabel("Avg Missed Precision-Recall Value")
plt.show()
```



Specify the minimum cross_validation error

```
In [38]: min_avg_missed_precision_recall_score = min(averaged_missed_precision_recall_scores)
print(min_avg_missed_precision_recall_score)

-0.03545454545454545
```